

```
In [1]: # Experiment No: 9
```

```
In [2]: # Aim: Simple Linear Regression
```

```
In [3]: # Name: Sakshi Padmakar Yeole
```

```
In [4]: # Class: 3rd year(B)
```

```
In [5]: # Roll No: 69
```

```
In [6]: # Date: 5th October 2024
```

```
In [7]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
```

```
In [8]: import pandas as pd
```

```
In [4]: import os
```

```
In [5]: os.getcwd()
```

```
Out[5]: 'C:\\Users\\hp'
```

```
In [6]: os.chdir("C:\\Users\\hp\\OneDrive\\Desktop")
```

```
In [9]: df=pd.read_csv("Salary.csv")
```

```
In [10]: df.head()
```

```
Out[10]:
```

| | YearsExperience | Salary |
|---|-----------------|--------|
| 0 | 1.1 | 39343 |
| 1 | 1.3 | 46205 |
| 2 | 1.5 | 37731 |
| 3 | 2.0 | 43525 |
| 4 | 2.2 | 39891 |

```
In [11]: df.tail()
```

```
Out[11]:
```

| | YearsExperience | Salary |
|----|-----------------|--------|
| 30 | 11.2 | 127345 |
| 31 | 11.5 | 126756 |
| 32 | 12.3 | 128765 |
| 33 | 12.9 | 135675 |
| 34 | 13.5 | 139465 |

```
In [12]: df.head(20)
```

```
Out[12]:
```

| | YearsExperience | Salary |
|---|-----------------|--------|
| 0 | 1.1 | 39343 |
| 1 | 1.3 | 46205 |

| | | |
|----|-----|-------|
| 2 | 1.5 | 37731 |
| 3 | 2.0 | 43525 |
| 4 | 2.2 | 39891 |
| 5 | 2.9 | 56642 |
| 6 | 3.0 | 60150 |
| 7 | 3.2 | 54445 |
| 8 | 3.2 | 64445 |
| 9 | 3.7 | 57189 |
| 10 | 3.9 | 63218 |
| 11 | 4.0 | 55794 |
| 12 | 4.0 | 56957 |
| 13 | 4.1 | 57081 |
| 14 | 4.5 | 61111 |
| 15 | 4.9 | 67938 |
| 16 | 5.1 | 66029 |
| 17 | 5.3 | 83088 |
| 18 | 5.9 | 81363 |
| 19 | 6.0 | 93940 |

In [13]: `df[5:10]`

Out[13]:

| | YearsExperience | Salary |
|---|-----------------|--------|
| 5 | 2.9 | 56642 |
| 6 | 3.0 | 60150 |
| 7 | 3.2 | 54445 |
| 8 | 3.2 | 64445 |
| 9 | 3.7 | 57189 |

In [14]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 35 entries, 0 to 34
Data columns (total 2 columns):
#   Column          Non-Null Count  Dtype
---  -
0   YearsExperience  35 non-null    float64
1   Salary          35 non-null    int64
dtypes: float64(1), int64(1)
memory usage: 688.0 bytes
```

In [15]: `df.describe()`

Out[15]:

| | YearsExperience | Salary |
|-------|-----------------|---------------|
| count | 35.000000 | 35.000000 |
| mean | 6.308571 | 83945.600000 |
| std | 3.618610 | 32162.673003 |
| min | 1.100000 | 37731.000000 |
| 25% | 3.450000 | 57019.000000 |
| 50% | 5.300000 | 81363.000000 |
| 75% | 9.250000 | 113223.500000 |
| max | 13.500000 | 139465.000000 |

In [16]: `df.shape`

Out[16]: (35, 2)

```
In [17]: df.size
```

```
Out[17]: 70
```

```
In [18]: df.ndim
```

```
Out[18]: 2
```

```
In [19]: df.columns
```

```
Out[19]: Index(['YearsExperience', 'Salary'], dtype='object')
```

```
In [20]: df.isnull().sum()
```

```
Out[20]: YearsExperience    0  
Salary                  0  
dtype: int64
```

```
In [21]: # Assigning values in X & Y  
x=df.iloc[:, :-1].values  
y=df.iloc[:, -1].values  
  
#x=df['year of experience']  
#y=df['salary']
```

```
In [22]: print(x)
```

```
[[ 1.1  
 [ 1.3]  
 [ 1.5]  
 [ 2. ]  
 [ 2.2]  
 [ 2.9]  
 [ 3. ]  
 [ 3.2]  
 [ 3.2]  
 [ 3.7]  
 [ 3.9]  
 [ 4. ]  
 [ 4. ]  
 [ 4.1]  
 [ 4.5]  
 [ 4.9]  
 [ 5.1]  
 [ 5.3]  
 [ 5.9]  
 [ 6. ]  
 [ 6.8]  
 [ 7.1]  
 [ 7.9]  
 [ 8.2]  
 [ 8.7]  
 [ 9. ]  
 [ 9.5]  
 [ 9.6]  
 [10.3]  
 [10.5]  
 [11.2]  
 [11.5]  
 [12.3]  
 [12.9]  
 [13.5]]
```

```
In [23]: print(y)
```

```
[ 39343  46205  37731  43525  39891  56642  60150  54445  64445  57189
```

```
63218 55794 56957 57081 61111 67938 66029 83088 81363 93940
91738 98273 101302 113812 109431 105582 116969 112635 122391 121872
127345 126756 128765 135675 139465]
```

```
In [24]: # Splitting testdata into x_train,x_test,y_train,y_test
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=.3,random_state=42)
```

```
In [25]: print(x_train)
```

```
[[12.9]
 [ 1.1]
 [ 2.2]
 [ 5.3]
 [ 9.6]
 [ 2.9]
 [ 4. ]
 [ 1.3]
 [ 1.5]
 [12.3]
 [ 2. ]
 [11.2]
 [ 8.2]
 [11.5]
 [ 3.9]
 [ 7.9]
 [ 5.9]
 [ 9. ]
 [ 3. ]
 [ 6.8]
 [13.5]
 [ 3.2]
 [ 4.5]
 [10.3]]
```

```
In [26]: print(x_test)
```

```
[[ 9.5]
 [ 4.1]
 [ 8.7]
 [ 7.1]
 [ 4.9]
 [10.5]
 [ 6. ]
 [ 4. ]
 [ 3.2]
 [ 5.1]
 [ 3.7]]
```

```
In [27]: print(y_train)
```

```
[135675 39343 39891 83088 112635 56642 55794 46205 37731 128765
 43525 127345 113812 126756 63218 101302 81363 105582 60150 91738
 139465 54445 61111 122391]
```

```
In [28]: print(y_test)
```

```
[116969 57081 109431 98273 67938 121872 93940 56957 64445 66029
 57189]
```

```
In [29]: from sklearn.linear_model import LinearRegression
lr=LinearRegression()
lr.fit(x_train,y_train)
```

```
Out[29]: LinearRegression()
```

```
In [30]: #Assigning Coefficient (slope) to m
m = lr.coef_
```

```
In [31]: print("Coefficient : " , m)
Coefficient : [8555.33918938]
```

```
In [32]: #Assigning Y-intercept to a
c = lr.intercept_
```

```
In [33]: print("Intercept : " , c)
Intercept : 29602.07353482097
```

```
In [34]: lr.score(x_test,y_test) * 100
Out[34]: 91.71426108885095
```

```
In [35]: df.isnull().any()
```

```
Out[35]: YearsExperience    False
Salary                    False
dtype: bool
```

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js